

Dantzig's Simplex Algorithm

How to Write Fast Numerical Code

by Rico Häuselmann
Donjan Rodic

Swiss Federal Institute of Technology (ETH Zurich)

27.05.2013



Linear Programming

Optimising a Linear Program in standard form:

Maximize

$$2x - 3y + z$$

Subject To

$$x + y + z \leq 10$$

$$4x - 3y + z \leq 3$$

$$2x + y - z \leq 6$$



Restrictions

- all coefficients positive (simplicity)
- all coefficients $\leq 10^6$ (stability)



Steps

- Tableau form

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 & 0 & 10 \\ 4 & -3 & 1 & 0 & 1 & 0 & 0 & 3 \\ 2 & 1 & -1 & 0 & 0 & 1 & 0 & 6 \\ 2 & -3 & 1 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

- Pivoting, reduced cost (objective function)
- Termination, worst runtime $O(e^m)$, but often $O(m)$



Steps

- Tableau form
$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 & 0 & 10 \\ 4 & -3 & 1 & 0 & 1 & 0 & 0 & 3 \\ 2 & 1 & -1 & 0 & 0 & 1 & 0 & 6 \\ 2 & -3 & 1 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$
- Pivoting, reduced cost (objective function)
- Termination, worst runtime $O(e^m)$, but often $O(m)$



Steps

- Tableau form
$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 & 0 & 10 \\ 4 & -3 & 1 & 0 & 1 & 0 & 0 & 3 \\ 2 & 1 & -1 & 0 & 0 & 1 & 0 & 6 \\ 2 & -3 & 1 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$
- Pivoting, reduced cost (objective function)
- Termination, worst runtime $O(e^m)$, but often $O(m)$



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Properties

- Tableau: $(m + 1) \times (m + n + 2)$
(requires full access each iteration)
Memory reads: $m(m + n) + 2m + n$
(all capacity misses for bigger problems)
Flops: $2 * m(m + n) + m$
- Computational intensity $I = \frac{2m^2 + 2mn + 4m + 2n}{8(m^2 + mn)} \sim \frac{1}{4}$



Properties

- Tableau: $(m + 1) \times (m + n + 2)$
(requires full access each iteration)
Memory reads: $m(m + n) + 2m + n$
(all capacity misses for bigger problems)
Flops: $2 * m(m + n) + m$
- Computational intensity $I = \frac{2m^2 + 2mn + 4m + 2n}{8(m^2 + mn)} \sim \frac{1}{4}$



Implementation

- **baseline**
- nta: cache control (inhibit polluting tableau data, prefetch next row)
- ssa: static scalar assignment
- block: reuse pivot row
- block_swap: swap pivot row at the end of tableau



Implementation

- baseline
- nta: cache control (inhibit polluting tableau data, prefetch next row)
- ssa: static scalar assignment
- block: reuse pivot row
- block_swap: swap pivot row at the end of tableau



Implementation

- baseline
- nta: cache control (inhibit polluting tableau data, prefetch next row)
- ssa: static scalar assignment
- block: reuse pivot row
- block_swap: swap pivot row at the end of tableau



Implementation

- baseline
- nta: cache control (inhibit polluting tableau data, prefetch next row)
- ssa: static scalar assignment
- block: reuse pivot row
- block_swap: swap pivot row at the end of tableau



Implementation

- baseline
- nta: cache control (inhibit polluting tableau data, prefetch next row)
- ssa: static scalar assignment
- block: reuse pivot row
- block_swap: swap pivot row at the end of tableau



Performance




Figure : performance comparison



Wall Time

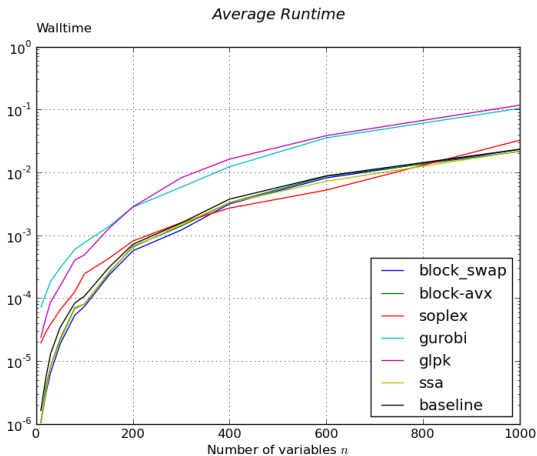


Figure : wall time comparison



Roofline

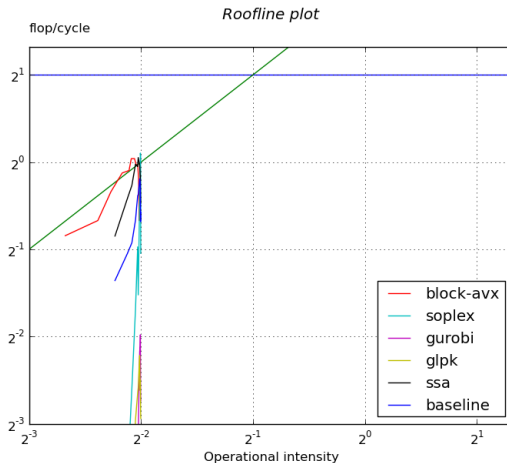


Figure : roofline



Profiling (gprof)

Memory load lines use most time

7.69	0.02	block_swap.hpp:122	T la1 = tabp[i*width+j];
3.85	0.01	block_swap.hpp:123	T la2 = tabp[i*width+j+1];
3.85	0.01	block_swap.hpp:124	T la3 = tabp[i*width+j+2];
3.85	0.01	block_swap.hpp:125	T la4 = tabp[i*width+j+3];
[... x16 ...]			

