

Dantzig's Simplex Algorithm

How to Write Fast Numerical Code

by Rico Häuselmann & Donjan Rodic

Swiss Federal Institute of Technology (ETH Zurich)

29.05.2013



Linear Programming

Optimising a Linear Program in standard form:

Maximize

$$2x - 3y + z$$

Subject To

$$x + y + z \leq 10$$

$$4x - 3y + z \leq 3$$

$$2x + y - z \leq 6$$



Restrictions

- all coefficients positive (simplicity)
- all coefficients $\leq 10^6$ (stability)



Steps

- Tableau form
$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 & 0 & 10 \\ 4 & -3 & 1 & 0 & 1 & 0 & 0 & 3 \\ 2 & 1 & -1 & 0 & 0 & 1 & 0 & 6 \\ 2 & -3 & 1 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$
- Pivoting, reduced cost (objective function), termination
- Worst runtime $O(e^m)$, but often $O(m)$



Steps

- Tableau form
$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 & 0 & 10 \\ 4 & -3 & 1 & 0 & 1 & 0 & 0 & 3 \\ 2 & 1 & -1 & 0 & 0 & 1 & 0 & 6 \\ 2 & -3 & 1 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$
- Pivoting, reduced cost (objective function), termination
- Worst runtime $O(e^m)$, but often $O(m)$



Steps

- Tableau form
$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 & 0 & 10 \\ 4 & -3 & 1 & 0 & 1 & 0 & 0 & 3 \\ 2 & 1 & -1 & 0 & 0 & 1 & 0 & 6 \\ 2 & -3 & 1 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$
- Pivoting, reduced cost (objective function), termination
- Worst runtime $O(e^m)$, but often $O(m)$



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Comparison

- GLPK (GNU Linear Programming Kit), solid standard solver
- CPLEX, mathematical OO-implementation
- Gurobi (CPLEX), fastest (multithreaded) solver available
- SoPlex, fastest FOSS solver available



Properties

- Tableau: $(m + 1) \times (m + n + 2)$
(requires full access each iteration)
Memory reads: $m(m + n) + 2m + n$
(all capacity misses for bigger problems)
Flops: $2 * m(m + n) + m$
- Computational intensity $I = \frac{2m^2 + 2mn + 4m + 2n}{8(m^2 + mn)} \sim \frac{1}{4}$



Properties

- Tableau: $(m + 1) \times (m + n + 2)$
(requires full access each iteration)
Memory reads: $m(m + n) + 2m + n$
(all capacity misses for bigger problems)
Flops: $2 * m(m + n) + m$
- Computational intensity $I = \frac{2m^2 + 2mn + 4m + 2n}{8(m^2 + mn)} \sim \frac{1}{4}$



Implementation

Elementary optimisations:

array: use raw pointers in place of `std::vector<std::vector>`

nta: cache control to inhibit polluting tableau data, prefetch rows

swap: store pivot row at a fixed location (end of tableau)

ssa: increase ILP via static assignment

blockX: reuse pivot row for X concurrent updates

sse/avx: use alignment & intrinsics to speed up float arithmetic



Framework

- Simplex runtime highly unpredictable
- Randomly generate LPs of increasing size (10-4000 vars)
- Grouped into 4 test sets

preview	(162 LPs, 145 MB)	quick testing
standard	(282 LPs, 260 MB)	plots
heavy	(852 LPs, 785 MB)	statistics
high	(50LPs, 3800 MB)	correctness & profiling



Framework

- Simplex runtime highly unpredictable
- Randomly generate LPs of increasing size (10-4000 vars)
- Grouped into 4 test sets

preview	(162 LPs, 145 MB)	quick testing
standard	(282 LPs, 260 MB)	plots
heavy	(852 LPs, 785 MB)	statistics
high	(50LPs, 3800 MB)	correctness & profiling



Framework

- Simplex runtime highly unpredictable
- Randomly generate LPs of increasing size (10-4000 vars)
- Grouped into 4 test sets

preview	(162 LPs, 145 MB)	quick testing
standard	(282 LPs, 260 MB)	plots
heavy	(852 LPs, 785 MB)	statistics
high	(50LPs, 3800 MB)	correctness & profiling



Framework

- Simplex runtime highly unpredictable
- Randomly generate LPs of increasing size (10-4000 vars)
- Grouped into 4 test sets

preview	(162 LPs, 145 MB)	quick testing
standard	(282 LPs, 260 MB)	plots
heavy	(852 LPs, 785 MB)	statistics
high	(50LPs, 3800 MB)	correctness & profiling



Performance

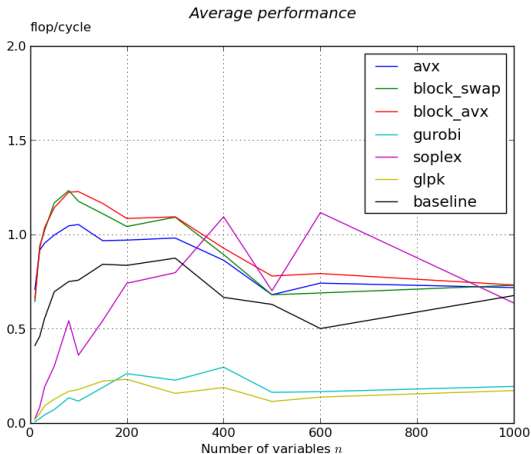


Figure: performance comparison



Wall Time

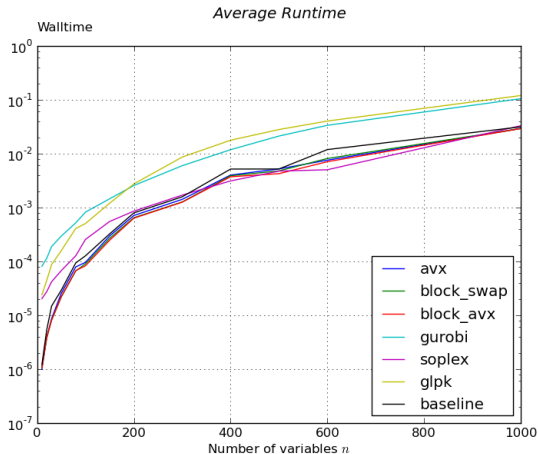


Figure: wall time comparison



Roofline

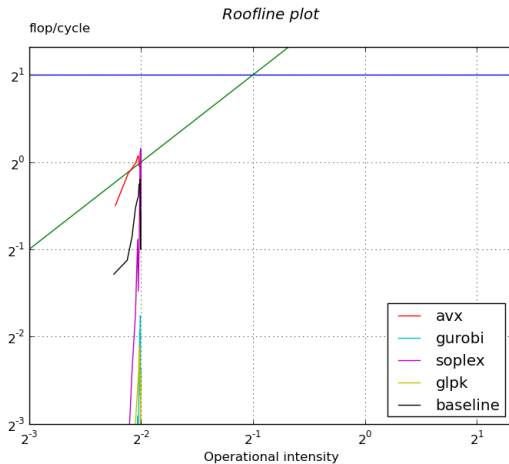


Figure: roofline



Profiling (gprof)

The main work routine has 24 memory access lines (12 shown):

%	sec	line
3.97	0.67	block_swap.hpp:122
4.81	0.81	block_swap.hpp:123
4.05	0.68	block_swap.hpp:124
4.53	0.76	block_swap.hpp:125
3.37	0.57	block_swap.hpp:128
4.00	0.67	block_swap.hpp:129
3.94	0.66	block_swap.hpp:130
4.47	0.75	block_swap.hpp:131
4.41	0.74	block_swap.hpp:139
5.07	0.85	block_swap.hpp:140
3.46	0.58	block_swap.hpp:141
2.98	0.50	block_swap.hpp:142

```
src/simplex/block_swap.hpp  
  
122  T r1 = tabp[m*width+j];  
123  T r2 = tabp[m*width+j+1];  
124  T r3 = tabp[m*width+j+2];  
125  T r4 = tabp[m*width+j+3];  
  
  
128  T la1 = tabp[i*width+j];  
129  T la2 = tabp[i*width+j+1];  
130  T la3 = tabp[i*width+j+2];  
131  T la4 = tabp[i*width+j+3];  
  
  
139  tabp[i*width+j]   = pa1;  
140  tabp[i*width+j+1] = pa2;  
141  tabp[i*width+j+2] = pa3;  
142  tabp[i*width+j+3] = pa4;
```



Profiling (valgrind)

Performance counters on a 1000 variables run:

	float add&mul	memory access
theoretical estimate	100'900'800	50'450'400
perf counters	100'825'000	50'704'015
perf counters for SSA	100'825'000	126'304'015
cachegrind profile	100'836'841	126'669'074



Profiling (perf)

Annotated perf recording on a 1000 variables run:

		T pa1 = la1 - fac1*r1;
		T pa2 = la2 - fac1*r2;
1.95	vmulsd %xmm5,%xmm0,%xmm10	
1.46	vmovsd (%rax),%xmm12	
		//~ PERFC_MEM += 4;
		T r1 = tabp[n*width+j];
		T r2 = tabp[n*width+j+1];
		T r3 = tabp[n*width+j+2];
		T r4 = tabp[n*width+j+3];
16.59	vmovsd 0x10(%rcx),%xmm3	

